

Message Oriented Middleware (MOM)

Message-based communication

Distributed System

Consists of a collection of distinct processes which:


- are spatially separated
- communicate with one another by exchanging messages

Message: atomic bit string; its format and meaning are specified by a communications protocol

Message-based Communication and Networking

The transport of a message from its source to its destination is performed by a computer network

The network can be abstracted as a communication channel



Internet Protocols

Layer	Function
Application	Specific communication services

Layer	Function
Transport	Communication between 2 (or more) processes
Network	Communication between 2 computers not directly connected with each other
Interface	Communication between 2 computers directly connected

On the Internet, the properties of the communication channel provided to an application depend on the transport protocol used (UDP/TCP):

- The design of a distributed application depends on the properties provided by the chosen transport protocol

UDP vs. TCP

Property	UDP	TCP
Abstraction	Message	Stream
Connection-based	✗	✓
Reliability (loss & duplication)	✗	✓
Order	✗	✓
Flow control	✗	✓
Number of recipients	1 / n	1

TCP Reliability

Message Loss

TCP guarantees that the application will be notified if the local end is unable to communicate with the remote end

- `send()/write()` or `recv()/read()` will return an error code
- TCP **cannot** guarantee that there is no data loss - the application should deal with this

Message duplication

Why not always retransmit?

The re-transmitted message may have been delivered before the connection was closed

- TCP is not able to filter data duplicated by the application

- Issue when it's a request for a non-idempotent operation (e.g. credit/debit operation)

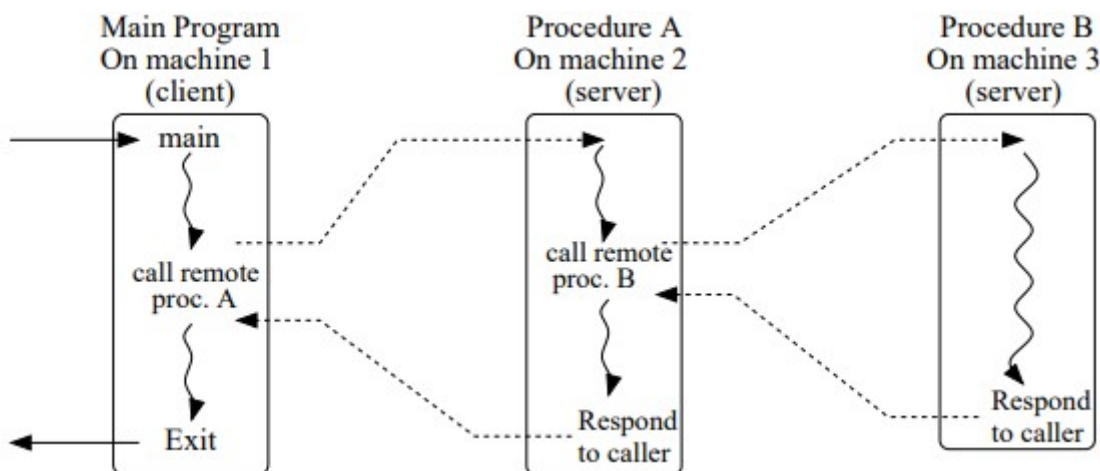
The application may need to **synchronize** with the remote end to learn if there was some data loss in either direction

RPC (Remote Procedure Call)

Reference: *van Steen and Tanenbaum, Distributed Systems, 3rd Ed.*

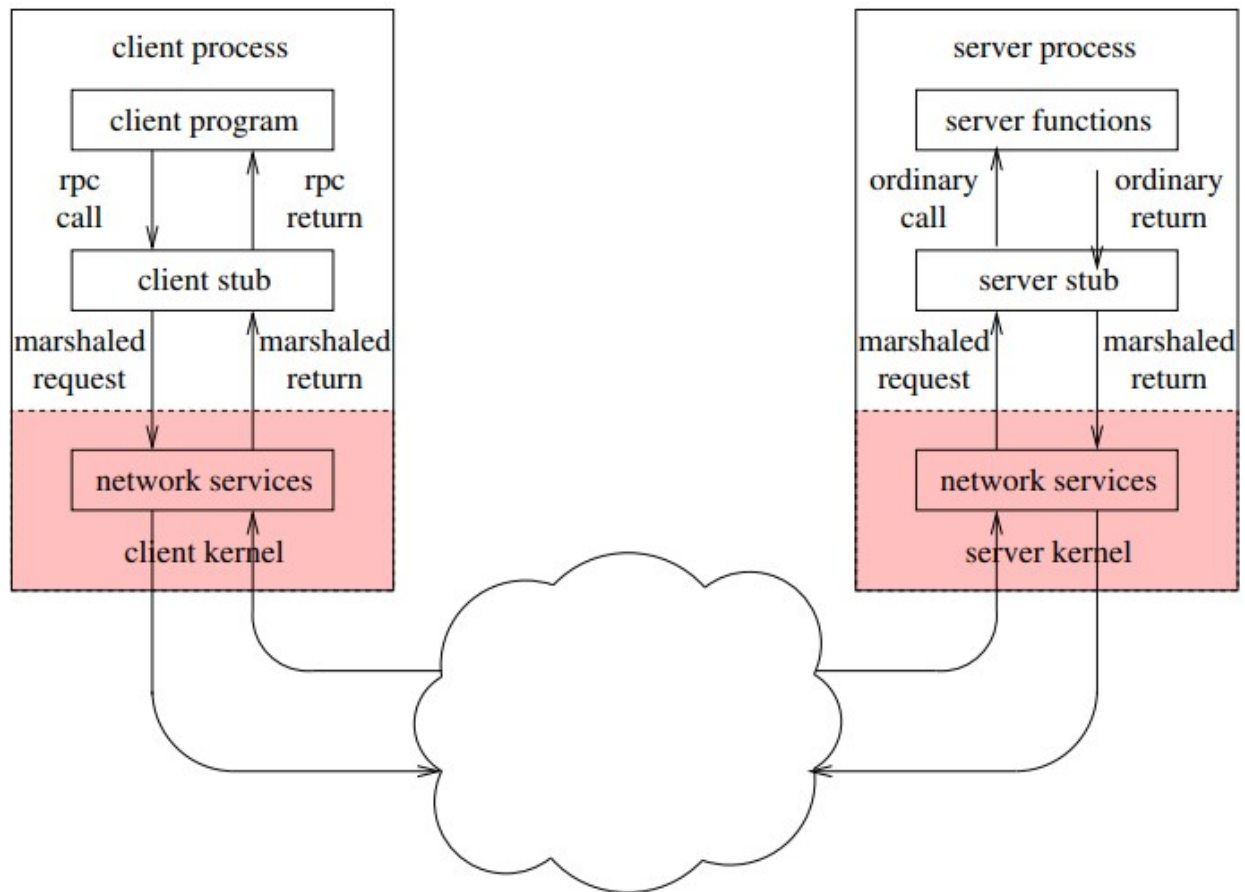
Idea: allow programs to call procedures located on other machines

1. Process on machine A calls procedure on machine B
 2. Calling process on A is suspended
 3. Execution of the called procedure takes place on B
- Information can be transported from the caller to the callee in the parameters and can come back in the procedure result.
 - No message passing at all is visible to the programmer.
 - Great for request-reply communication patterns (but even then there may be better alternatives)



Architecture

- RPC is typically implemented on top of the transport layer (TCP/IP)



Client Stub

Request

1. Assembles message: **parameter marshalling**
2. Send messages (`write()`/`sendto()`)
3. Blocks waiting for response (`read()`/`recvfrom()`)
 - Assuming synchronous RPC

Response

1. Receives response
2. Extracts the results (**unmarshalling**)
3. Returns to client
 - Assuming synchronous RPC

Server Stub

Request

1. Receives message (`read()`/`recvfrom()`)
2. Parses message to determine arguments (**unmarshalling**)
3. Calls function

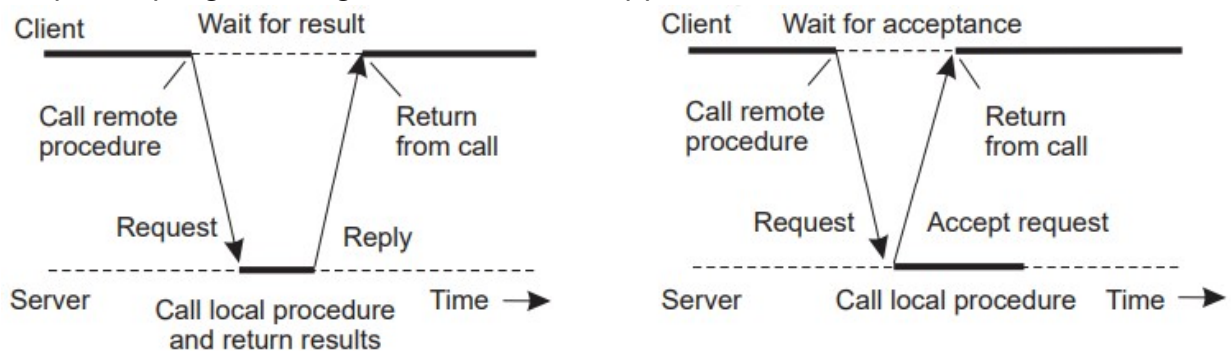
Response

1. Assembles message with the return value of the function
2. Send message (`write()`/`sendto()`)
3. Blocks waiting for a new request

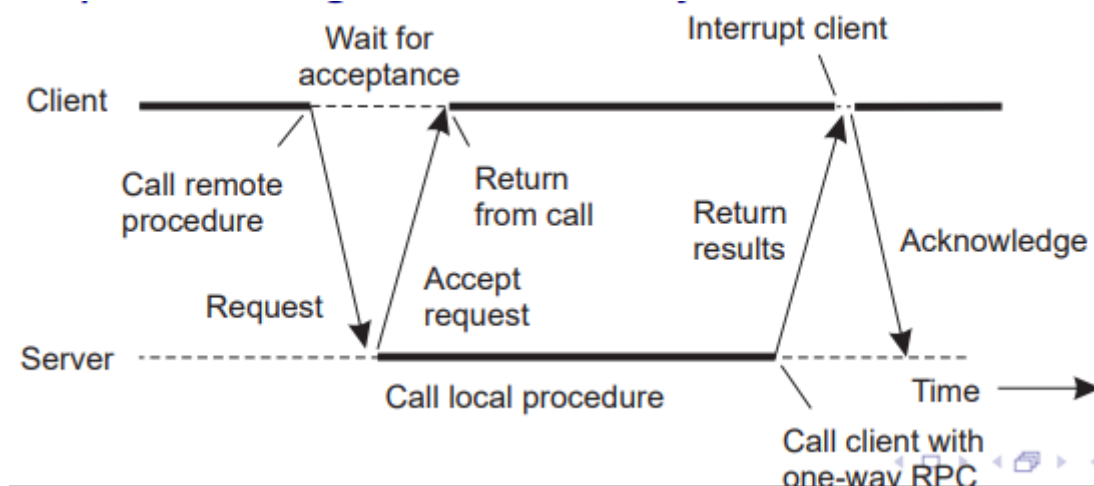
Synchronous vs. Asynchronous

Synchronous

- Client blocks until it receives the response
- Programming distributed applications with synchronous RPC would be almost as simple as programming non-distributed applications



Asynchronous



Asynchronous Communication (MOM)

Problem: The communication parties may not always be simultaneously available

Solution: Use *asynchronous communication*

Message Oriented Middleware (MOM)

Asynchronous message-based communication

- Communication service (middleware) stores the messages as long as needed to deliver them
 - Guarantees **asynchronicity** and **anonymity**
- Publishers:** may send messages
Subscribers: may receive messages

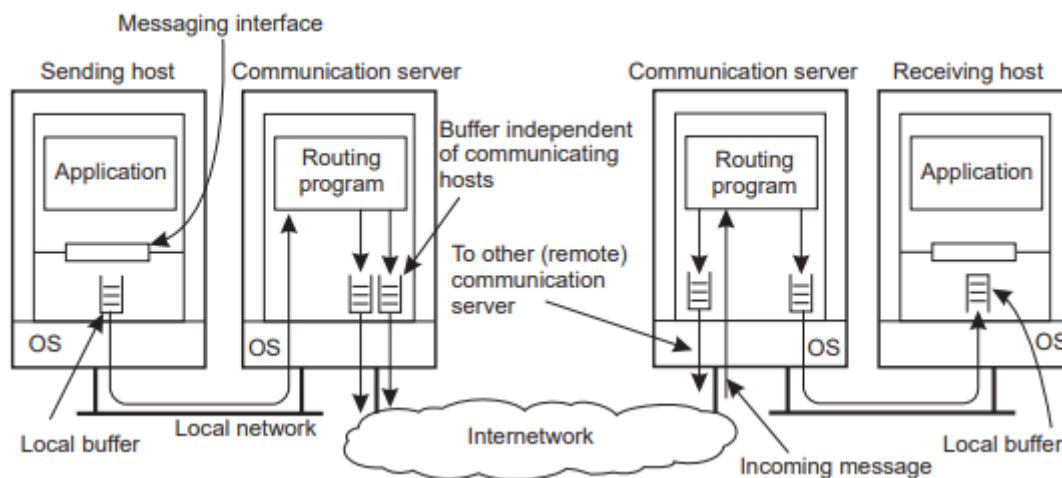
Basic Patterns

Pattern	Structure	Roles	Delivery
Point-to-point	Queue	Senders and Receivers	Each message is delivered to at most one process
Publish-Subscribe	Topic	Publishers and Subscribers	A message is delivered to more than one process

Messaging Service Implementation

Provides asynchronous communication

- At the lowest communication level, there must be synchronization between sender and receiver



Asynchronous Communication Applications

Appropriate for applications when the send and receiver are **loosely coupled**:

- Enterprise Application Integration
- Workflow Applications
 - Related to **business processes**
- Microservices
- Message based communication between people (Email, SMS, Instant Messaging)

Java Message Service (JMS)

Temos se saber cenas especificas de java?.....

TODO

Message Queing Protocols

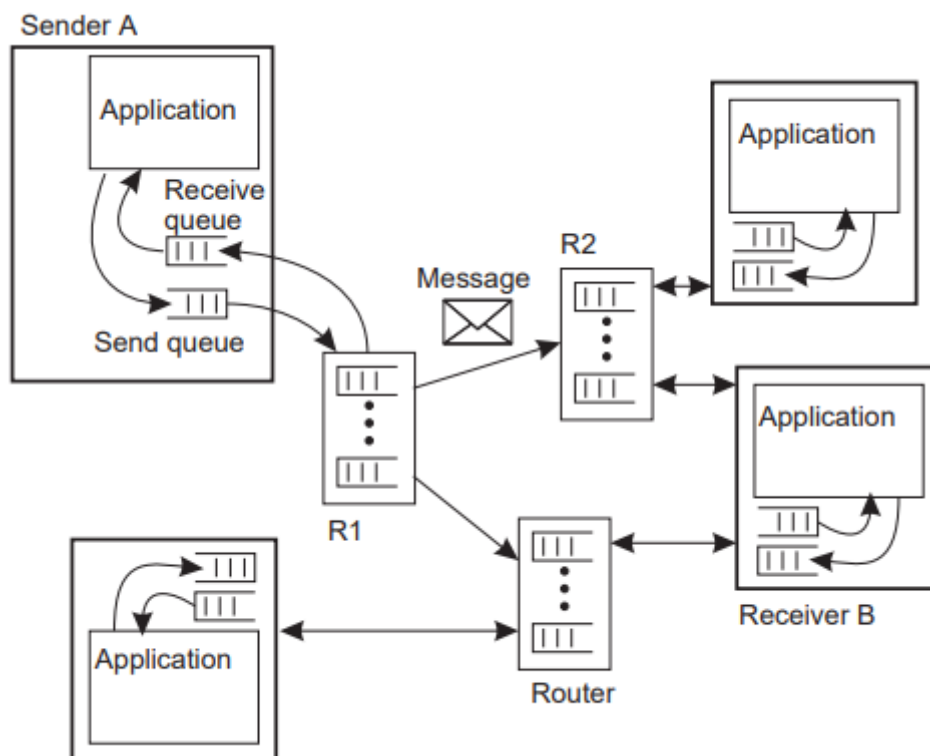
AMQP (Advanced Message Queuing Protocol): open-standard protocol (approved by OASIS)

MTTQ (Message Queuing Telemetry Transport): also an OASIS transport

OpenWire: public protocol used by Apache ActiveMQ

Architecture

- Larger scale systems may use **message relays** to route messages to their destinations
- Similar to SMTP, although nowadays almost every email message just traverses 2 servers



Message Brokers

Convert the format of the messages used by one application to the format used by another application

- They are not part of the communication service