

Replication and Consistency Models

Data Replication

Replicate data at many nodes

- **Performance:** local reads
- **Reliability:** no data-loss unless data is lost in all replicas
- **Availability:** data available unless all replicas fail or become unreachable
- **Scalability:** balance load across nodes for reads

Upon an update, push data to all replicas


Challenge: Ensure data consistency

- Updating at different replicas may lead to different results e.g. inconsistent data

Strong Consistency

All replicas execute updates in the same order

- same initial state leads to same result




Models

Sequential Consistency

An execution is sequential consistent iff it is identical to a sequential execution of all the operations in that execution such that all operations executed by any thread, appear in the order in which they were executed by the corresponding thread.

- Model provided by a multi-threaded system on a uniprocessor
- Protocol
 - **Read:** reads from one replica

- **Write:** writes to all replicas in same order; have no reply → return after sending the write request messages to all replicas
- **Snapshot:** reads from one replica
- Is not Composable
 - the same algorithm to replicate each of the sub-arrays, and thus ensure sequential consistency on each array
 - The combined execution may not be sequential consistent



1. *write(1, 5)*

2. *read(1)*

3. *read(2)*

4. *write(2, 7)*

5. *snapshot()*

6. *write(3, 2)*


What other ordering would be possible?

Linearizability

An execution is linearizable iff it is sequential consistent and if op_1 occurs before op_2 , according to one omniscient observer, then op_1 appears before op_2 .

- Guaranteeing linearizability usually requires more synchronization
- Protocol
 - **Read:** reads from one replica
 - **Write:** writes to all replicas in same order; **Wait for ack from all replicas before returning**

- **Snapshot:** reads from one replica



src: Fekete and Ramamritham 09

1. *write(1, 5)*

2. *read(1)*

3. *read(2)*

4. *write(2, 7)*

5. *snapshot()*

6. *write(3, 2)*

Is the other sequentially consistent order also linearizable?

One-copy Serializability (Transaction-based Systems)

The execution of a set of transactions is one-copy serializable iff its outcome is similar to the execution of those transactions in a single copy

- it's essentially the sequential consistency model, when the operations executed by all processor are transactions
 - isolation: ensures that the outcome of the concurrent execution of a set of transactions is equal to some sequential execution of those transactions
- **Serializability:** databases (preserve complex application-specific invariants)
- **Sequential Consistency:** multiprocessing (programmers are expected to reason about concurrency)