

# Replication for Fault Tolerance

## Quorum Consensus Protocols

- Clients communicate directly to the servers/replicas
- Each (replicated) operation (e.g. read/write) requires a **quorum** (set of replicas)

### Fundamental Property:

- If the result of one operation depends on the result of another, then their quorums must overlap, i.e. have common replicas

### Defining Quorums

- Consider all replicas as peers
  - quorums are determined by their size, i.e. the number of replicas in the quorum
  - This is equivalent to assign 1 vote to each replica

## Read/Write Quorums Must Overlap


- The replicas provide **only read** and **write** operations
- Because the output of a read operation depends on previous write operations, the read quorum must overlap the write quorum:

**$N_R + N_W > N$** , where

$N_R$  → size of the read quorum

$N_W$  → size of the write quorum

$N$  → number of replicas



## Implementation

Each object's replica has a version number

- **Read**

1. Poll a read quorum, to find out the current version
  - A server replies with the current version
2. Read the object value from an up-to-date replica
  - If the size of the object is small, it can be read as the read quorum is assembled

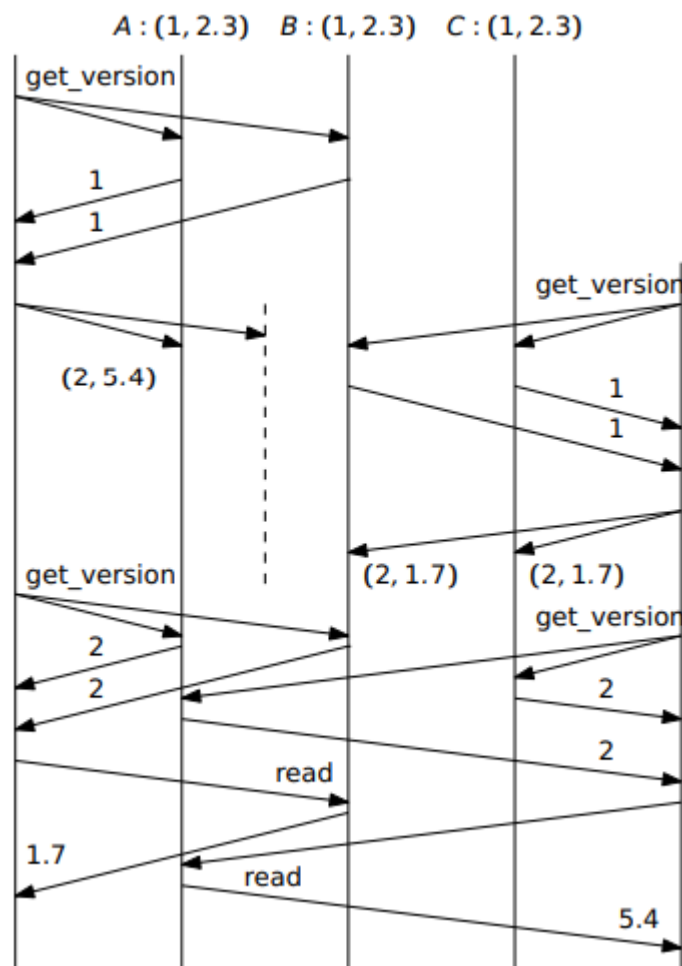
- **Write**

1. Poll a write quorum, to find out the current version
  - A server replies with the current version
2. Write the new value with the new version to a write quorum
  - We assume that writes modify the entire object, not parts of it

Note: A write operation depends on previous write operations (via the version) and therefore write quorums must overlap:  $NW + NW > N$

## Naive Implementation with Faults

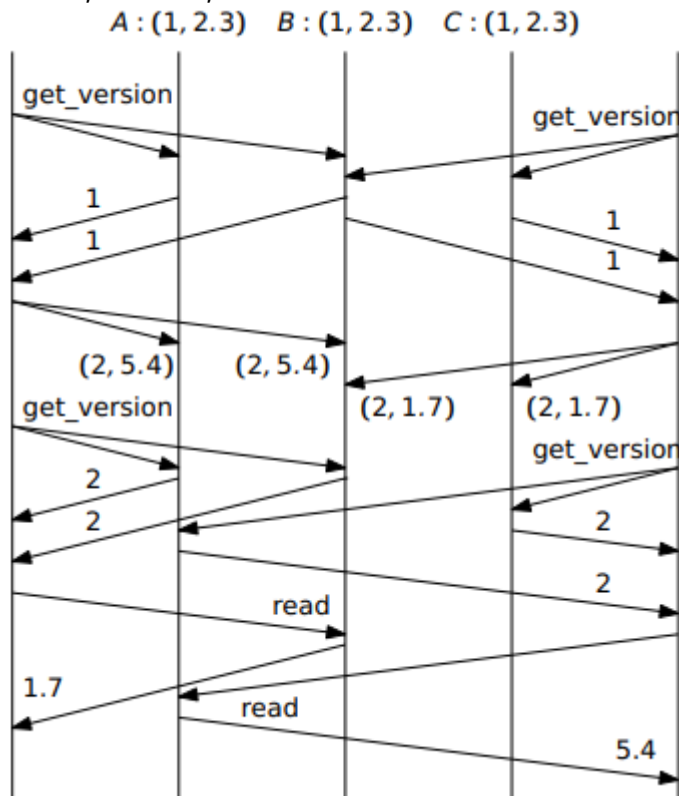
- $N = 3$ ,  $NR = 2$ ,  $NW = 2$



- First/left client attempts to write, but because of a partition it updates only one replica (A)
- Second/right client, in different partition, attempts to write and it succeeds.
- Variable has different values for the same version.
- The partition heals and each client does a read
- Each client gets a value different from the one it wrote.

## Naive Implementation with Concurrent Writes

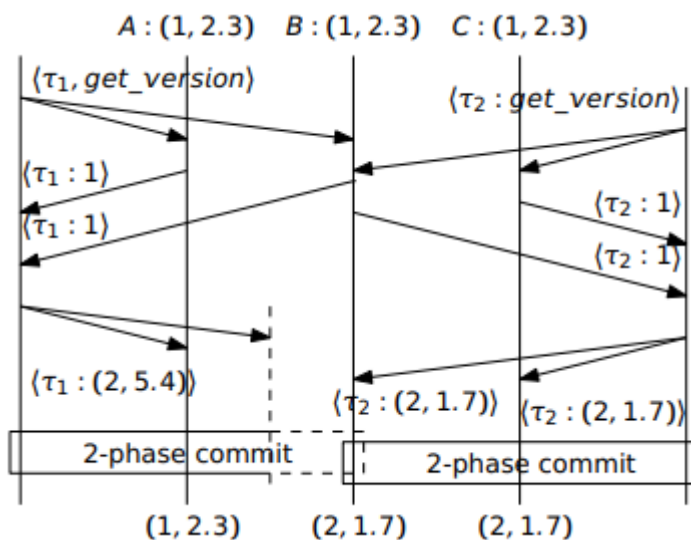
- $N = 3$ ,  $NR = 2$ ,  $NW = 2$



- Two clients attempt to write the replicas at more or less the same time
- The two write quorums are not equal, even though they overlap
- Again, replicas end up in an inconsistent state.
- Soon after, each client does a read
- Each client gets a value different from the one it wrote.

## Ensuring Consistency with Transactions

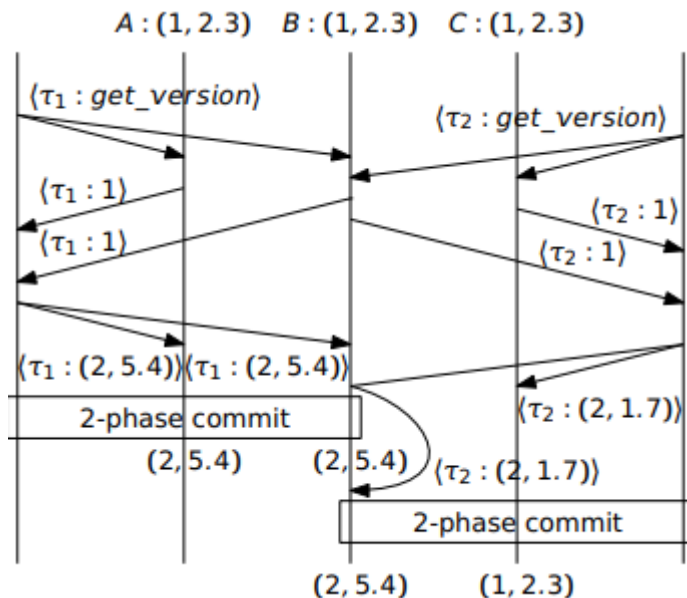
- The write (or read) of each replica is an operation of a distributed transaction
- If the write is not accepted by at least a write quorum, the transaction aborts



- The left client will not get the vote from replica B and therefore it will abort transaction  $\tau_1$
- On the other hand, transaction  $\tau_2$  commits, and its write will be effective.

Transactions also prevent inconsistencies in the case of concurrent writes

- Transactions ensure isolation, by using concurrency control
- Lets assume the use of locks.



- Server B processes the left client write request first, and acquires a write lock on behalf of  $\tau_1$
- When server B processes the right client write request, it tries to acquire a write lock on behalf of  $\tau_2$ , but it is forced to wait for the termination of  $\tau_1$
- The commit of  $\tau_1$  in server B invalidates the version number of  $\tau_2$ 's write and therefore  $\tau_2$  aborts.

## XA-based Quorum Consensus Implementation

Each object's access is performed in the context of a transaction

### Read

1. Poll a read quorum, to find out the current version
  - There is no need to read the object's state
  - Only the first time the transaction reads the object
2. Read the object state from an up-to-date replica
  - Only the first time the transaction reads the project

### Write

1. Poll a write quorum, to find out the current version and which replicas are up-to-date

- On the first time the transaction writes the object:
  - Object state may have to be read from an up-to-date replica
  - Replicas may have to be updated
- 2. Write the new value with the new version
  - Replica rejects write if version is **not valid**
  - All writes by a transaction are applied to the same replicas
    - Because these will be the only ones with an up-to-date version

## Transaction-based Quorum Consensus Replication

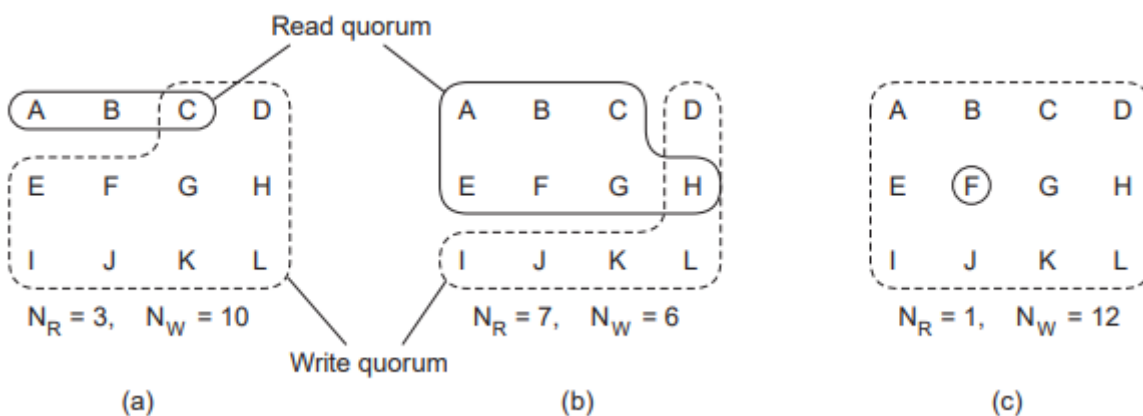
### Pros

- Transactions solve both the problem of failures and concurrency
- Transactions can also support a more complex computations

### Cons

- Deadlocks are possible, if transactions use locks
- Blocking if transactions use two-phase commit
- Availability problems

## Playing with Quorums



(c) → **Read-one/Write-all protocol**

- By assigning each replica its own number of votes, which may be different from one, **weighted-voting** provides extra flexibility.

## Quorum Consensus Fault Tolerance

- tolerates unavailability of replicas
- The availability analysis by Gifford relies on the probability of crashing of a replica/server

# Dynamo Quorums

## Dynamo

- replicated key-value storage system developed at Amazon
- uses quorums to provide high-availability
- enhances high-availability, by using multi-version objects

## Dynamo's Quorums

- Each key is associated with a set of servers, the **preference list**
  - The first N servers in this list are the main replicas
  - The remaining servers are backup replicas and are used only in the case of failures
- Each operation (get()/put()) has a coordinator, which is one of the first N servers in the preference list.
  - The coordinator is the process that executes the actions typically executed by the client in Gifford's quorums

put(.) → requires a quorum of W replicas

get(.) → requires a quorum of R replicas

such that:

$$R + W > N$$

### put(key, value, context) coordinator

1. Generates the version vector for the new version and writes the new value locally
  - The new version vector is determined by the coordinator from the **context**, a set of version vectors
2. Send the (key, value) and its version vector to the N first servers in the key's preference list
  - The put() is deemed successful if at least W-1 replicas respond

### get(key) coordinator

1. Requests all versions of the (key, value) pair, including the respective version vectors, from the remaining first N servers in the preference list
2. On receiving the response from at least R-1 replicas, it returns all the (key,value) pairs whose version-vector are **maximal**
  - If there are multiple pairs, the application that executed the get() is supposed to reconcile the different versions and write-back the reconciled pair using put().

## Dynamo's "Sloppy" Quorums and Hinted Handoff

- **Without failures** Dynamo provides strong consistency
- **In the case of failures** the coordinator may not be able to get a quorum from the N first replicas in the preference list

To **ensure availability**:

- the coordinator will try to get a **sloppy quorum** by enlisting the backup replicas in the preference list
  - The copy of the (key, value) sent to the backup server has a **hint** in its metadata identifying the server that was supposed to keep that copy
  - The backup server scans periodically the servers it is substituting

**At the cost of consistency** sloppy quorums do not ensure that every quorum of a get() overlaps every quorum of a put()

- **Sloppy quorums** are intended as a solution to temporary failures
  - To handle failures with a longer duration, Dynamo uses a anti-entropy approach for replica synchronization